

Python Database Access

John Strickler

Version 1.0, March 2021

Table of Contents

Chapter 1: Database Access	1
The DB API	2
Connecting to a Server	3
Creating a Cursor	5
Executing a Statement	6
Fetching Data	7
SQL Injection	10
Parameterized Statements	12
Dictionary Cursors	19
Metadata	23
Transactions	26
Object-relational Mappers	27
NoSQL	28
Index	35

Chapter 1: Database Access

Objectives

- Understand the Python DB API architecture
- Connect to a database
- Execute simple and parameterized queries
- Fetch single and multiple row results
- Get metadata about a query
- Execute non-query statements
- Start transactions and commit or rollback as needed

The DB API

- Several ways to access DBMSs from Python
- DB API is most popular
- DB API is sort of an "abstract class"
- Many modules for different DBMSs using DB API
- Hides actual DBMS implementation

To make database programming simpler, Python has the DB API. This is an API to standardize working with databases. When a package is written to access a database, it is written to conform to the API, and thus programmers do not have to learn a new set of methods and functions.

Table 1. Available Interfaces (using Python DB API-2.0)

Database	Python package
Firebird (and Interbase)	KInterbasDB
IBM DB2	PyDB2
Informix	informixdb
Ingres	ingmod
Microsoft SQL Server	pymssql
MySQL	pymysql
ODBC	pyodbc
Oracle	cx_oracle
PostgreSQL	psycopg2
SAP DB (also known as "MaxDB")	sapdbapi
SQLite	sqlite3
Sybase	Sybase

NOTE There may be other interfaces to some of the listed DBMSs as well.

Connecting to a Server

- Import appropriate library
- Use `connect()` to get a database object
- Specify host, database, username, password

To connect to a database server, import the package for the specific database. Use the package's **connect()** method to get a database object, specifying the host, initial database, username, and password. If the username and password are not needed, use `None`.

Some database modules have nonstandard parameters to the `connect()` method.

When finished with the connection, call the **close()** method on the connection object.

Many database modules support the context manager (**with** statement), and will automatically close the database when the `with` block is exited. Check the documentation to see how this is implemented for a specific database.

Example

```
import sqlite3

slconn = sqlite3.connect('web_content')

import pymysql

myconn = pymysql.connect (host = "myserver1",
                          user = "adeveloper",
                          passwd = "s3cr3t",
                          db = "web_content")

# make queries, etc. here ...
myconn.close()
```

NOTE

Argument names for the `connect()` method may not be consistent. For instance, `pymysql` supports the above parameter names, while `pymssql` does not.

Table 2. `connect()` examples

Package	Database	Connection
cx_oracle	Oracle	<pre>ip = 'localhost' port = 1521 SID = 'YOURSIDHERE' dsn_tns = cx_Oracle.makedsn(ip, port, SID) db = cx_Oracle.connect('adeveloper', '\$3cr3t', dsn_tns)</pre>
psychopg	PostgreSQL	<pre>psycpg2.connect ('' host='localhost' user='adeveloper' password='\$3cr3t' dbname='testdb' '')</pre> <p><i>note: connect() has one (string) parameter, not multiple parameters</i></p>
pymssql	MS-SQL	<pre>pymssql.connect (host="localhost", user="adeveloper", passwd="\$3cr3t", db="testdb",)</pre> <pre>pymssql.connect (dsn="DSN",)</pre>
pymysql	MySQL	<pre>pymysql.connect (host="localhost", user="adeveloper", passwd="\$3cr3t", db="testdb",)</pre>
pyodbc	Any ODBC-compliant DB	<pre>pyodbc.connect('' DRIVER={SQL Server}; SERVER=localhost; DATABASE=testdb; UID=adeveloper; PWD=\$3cr3t '')</pre> <pre>pyodbc.connect('DSN=testdsn;PWD=\$3cr3t')</pre> <p><i>note: connect() has one (string) parameter, not multiple parameters</i></p>
sqlite3	SQLite3	<pre>sqlite3.connect('testdb')</pre> <pre>sqlite3.connect(':memory:')</pre>

Creating a Cursor

- Cursor can execute SQL statements
- Multiple cursors available
 - Standard cursor
 - Returns tuples
 - Other cursors
 - Returns dictionaries
 - Leaves data on server

Once you have a database object, you can create one or more cursors. A cursor is an object that can execute SQL code and fetch results.

The default cursor for most packages returns each row as a tuple of values. There are different types of cursors that can return data in different formats, or that control whether data is stored on the client or the server.

Example

```
myconn = pymysqlconnect (host="myserver1",user="adeveloper", passwd="s3cr3t",  
db="web_content")  
mycursor = myconn.cursor()
```

Executing a Statement

- Executing cursor sends SQL to server
- Data not returned until asked for
- Returns number of lines in result set for queries
- Returns lines affected for other statements

Once you have a cursor, you can use it to perform queries, or to execute arbitrary SQL statements via the `execute()` method. The first argument to `execute()` is a string containing the SQL statement to run.

Example

```
cursor.execute("select hostname,ostype,user from hostinfo")
cursor.execute('insert into hostinfo values
("foo",5,"2.6","arch","net",2055,3072,"bob",0)')
```

Fetching Data

- Use one of the fetch methods from the cursor object
- Syntax
 - `rec = cursor.fetchone()`
 - `recs = cursor.fetchall()`
 - `recs = cursor.fetchmany()`

Cursors provide three methods for returning query results.

fetchone() returns the next available row from the query results.

fetchall() returns a tuple of all rows.

fetchmany(n) returns up to n rows. This is useful when the query returns a large number of rows.

Example

```
cursor.execute("select color, quest from knights where name = 'Robin'")
(color,quest) = cursor.fetchone()

cursor.execute("select color, quest from knights")
rows = cursor.fetchall()

cursor.execute("select * from huge_table")
while True:
    rows = cursor.fetchmany(1000)
    if not rows:
        break
    for row in rows:
        # process row
```

Example

db_sqlite_basics.py

```
#!/usr/bin/env python

import sqlite3

with sqlite3.connect("../DATA/presidents.db") as s3conn: ①

    s3cursor = s3conn.cursor() ②

    # select first name, last name from all presidents
    s3cursor.execute('''
        select firstname, lastname
        from presidents
    ''') ③

    print("Sqlite3 does not provide a row count\n") ④

    for row in s3cursor.fetchall(): ⑤
        print(' '.join(row)) ⑥
```

- ① connect to the database
- ② get a cursor object
- ③ execute a SQL statement
- ④ (included for consistency with other DBMS modules)
- ⑤ fetchall() returns all rows
- ⑥ each row is a tuple

db_sqlite_basics.py

```
Richard Milhous Nixon
Gerald Rudolph Ford
James Earl 'Jimmy' Carter
Ronald Wilson Reagan
George Herbert Walker Bush
William Jefferson 'Bill' Clinton
George Walker Bush
Barack Hussein Obama
Donald J Trump
Joseph Robinette Biden
```

NOTE

See `db_mysql_basics.py` and `db_postgres_basics.py` for examples using those modules. In general, all the `sqlite3` examples are also implemented for MySQL and Postgres, plus a few extras.

SQL Injection

- "Hijacks" SQL code
- Result of string formatting
- Always use parameterized statements

One kind of vulnerability in SQL code is called SQL injection. This occurs when an attacker embeds SQL commands in input data. This can happen when naively using string formatting to build SQL statements:

Example

db_sql_injection.py

```
#!/usr/bin/env python
#
good_input = 'Google'
malicious_input = "'; drop table customers; -- " ①

naive_format = "select * from customers where company_name = '{} ' and company_id != 0"

good_query = naive_format.format(good_input) ②
malicious_query = naive_format.format(malicious_input) ②

print("Good query:")
print(good_query) ③
print()

print("Bad query:")
print(malicious_query) ④
```

- ① input would come from a web form, for instance
- ② string formatting naively adds the user input to a field, expecting only a customer name
- ③ non-malicious input works fine
- ④ query now drops a table ('--' is SQL comment)

db_sql_injection.py

Good query:

```
select * from customers where company_name = 'Google' and company_id != 0
```

Bad query:

```
select * from customers where company_name = ''; drop table customers; -- ' and  
company_id != 0
```

Parameterized Statements

- More efficient updates
- Use placeholders in query
 - Placeholders vary by DB
- Pass iterable of parameters
- Prevent SQL injection
- Use `cursor.execute()` or `cursor.executemany()`

For efficiency, you can iterate over a sequence of input datasets when performing a non-query SQL statement. The `execute()` method takes a query, plus an iterable of values to fill in the placeholders. The database manager will only parse the query once, then reuse it for subsequent calls to `execute()`.

Parameterized queries also protect against SQL injection attacks.

Different database modules use different placeholders. To see what kind of placeholder a module uses, check `MODULE.paramstyle`. Types include `'pyformat'`, meaning `'%s'`, and `'qmark'`, meaning `'?'`.

The `executemany()` method takes a query, plus an iterable of iterables. It will call `execute()` once for each nested iterable.

Example

```
single_row = ("Smith","John","green"),

multi_rows= [
    ("Smith","John","green"),
    ("Douglas","Sam","pink"),
    ("Robinson","Alberta","blue"),
]

query = "insert into people (lname,fname,color) values (%s,%s,%s)"

rows_added = cursor.execute(query, single_row)
rows_added = cursor.executemany(query, multi_rows)
```

Table 3. Placeholders for SQL Parameters

Python package	Placeholder for parameters
pymysql	%s
cx_oracle	:param_name
pyodbc	?
pymssql	%d for int, %s for str, etc.
Psychopg	%s or %(param_name)s
sqlite3	? or :param_name

TIP with the exception of **pymssql** the same placeholder is used for all column types.

Example

db_sqlite_parameterized.py

```
#!/usr/bin/env python

import sqlite3

with sqlite3.connect("../DATA/presidents.db") as s3conn:
    s3cursor = s3conn.cursor()

    party_query = '''
    select firstname, lastname
    from presidents
    where party = ?
    ''' ①

    for party in 'Federalist', 'Whig':
        print(party)
        s3cursor.execute(party_query, (party,)) ②
        print(s3cursor.fetchall())
        print()
```

① ? is SQLite3 placeholder for SQL statement parameter; different DBMSs use different placeholders

② second argument to execute() is iterable of values to fill in placeholders from left to right

db_sqlite_parameterized.py

```
Federalist
[('John', 'Adams')]

Whig
[('William Henry', 'Harrison'), ('John', 'Tyler'), ('Zachary', 'Taylor'), ('Millard',
'Fillmore')]
```

Example

db_sqlite_bulk_insert.py

```
#!/usr/bin/env python
import os
import sqlite3
import random

FRUITS = ["pomegranate", "cherry", "apricot", "date", "apple",
          "lemon", "kiwi", "orange", "lime", "watermelon", "guava",
          "papaya", "fig", "pear", "banana", "tamarind", "persimmon",
          "elderberry", "peach", "blueberry", "lychee", "grape"]

DB_NAME = 'fruitprices.db' ①

CREATE_TABLE = """
create table fruit (
    name varchar(30),
    price decimal
)
""" ②

INSERT = '''
insert into fruit (name, price) values (?, ?)
''' ③

def main():
    """
    Program entry point.

    :return: None
    """
    conn = get_connection()
    create_database(conn)
    populate_database(conn)

    read_database()

def get_connection():
    """
    Get a connection to the PRODUCE database

    :return: SQLite3 connection object.
    """
    if os.path.exists(DB_NAME):
```

```
os.remove(DB_NAME) ④

s3conn = sqlite3.connect(DB_NAME) ⑤
return s3conn

def create_database(conn):
    """
    Create the fruit table

    :param conn: The database connection
    :return: None
    """
    conn.execute(CREATE_TABLE) ⑥

def populate_database(conn):
    """
    Add rows to the fruit table

    :param conn: The database connection
    :return: None
    """

    fruit_data = get_fruit_data() # [('apple', .49), ('kiwi', .38)]

    try:
        conn.executemany(INSERT, fruit_data) ⑦
    except sqlite3.DatabaseError as err:
        print(err)
        conn.rollback()
    else:
        conn.commit() ⑧

def get_fruit_data():
    """
    Create iterable of fruit records.

    :return: Generator of name/price tuples.
    """
    return ((f, round(random.random() * 10 + 5, 2)) for f in FRUITS) ⑨

def read_database():
    conn = sqlite3.connect(DB_NAME)
    for name, price in conn.execute('select name, price from fruit'):
        print('{:12s} {:6.2f}'.format(name, price))
```

```
if __name__ == '__main__':  
    main()
```

- ① set name of database
- ② SQL statement to create table
- ③ parameterized SQL statement to insert one record
- ④ remove existing database if it exists
- ⑤ connect to (new) database
- ⑥ run SQL to create table
- ⑦ iterate over list of pairs and add each pair to the database
- ⑧ commit the inserts; without this, no data would be saved
- ⑨ build list of tuples containing fruit, price pairs

db_sqlite_bulk_insert.py

pomegranate	12.85
cherry	8.66
apricot	10.49
date	8.19
apple	5.44
lemon	5.53
kiwi	9.58
orange	8.50
lime	10.01
watermelon	10.45
guava	12.68
papaya	12.92
fig	12.37
pear	11.82
banana	13.42
tamarind	8.01
persimmon	9.53
elderberry	5.72
peach	13.19
blueberry	10.10
lychee	12.75
grape	8.62

Dictionary Cursors

- Indexed by column name
- Not standardized in the DB API

The standard cursor provided by the DB API returns a tuple for each row. Most DB packages provide other kinds of cursors, including user-defined versions.

A very common cursor is a dictionary cursor, which returns a dictionary for each row, where the keys are the column names. Each package that provides a dictionary cursor has its own way of providing the dictionary cursor, although they all work the same way.

For the packages that don't have a dictionary cursor, you can make a generator function that will emulate one.

Table 4. Dictionary Cursors

Python package	How to get a dictionary cursor
pymysql	<pre>import pymysql.cursors conn = pymysql.connect(..., cursorclass = pymysql.cursors.DictCursor) dcur = conn.cursor() <i>all cursors will be dict cursors</i></pre> <pre>dcur = conn.cursor(pymysql.cursors.DictCursor) <i>only this cursor will be a dict cursor</i></pre>
cx_oracle	<i>Not available</i>
pyodbc	<i>Not available</i>
pgdb	<i>Not available</i>
pymssql	<pre>conn = pymssql.connect (... , as_dict=True) dcur = conn.cursor()</pre>
psycopg	<pre>import psycopg2.extras dcur = conn.cursor(cursor_factory=psycopg.extras.DictCu rsor)</pre>
sqlite3	<pre>conn = sqlite3.connect (... , row_factory=sqlite3.Row) dcur = conn.cursor() conn.row_factory = sqlite3.Row dcur = conn.cursor()</pre>

Example

db_sqlite_extras.py

```
#!/usr/bin/env python

import sqlite3

s3conn = sqlite3.connect("../DATA/presidents.db")
# uncomment to make _all_ cursors dictionary cursors
# conn.row_factory = sqlite3.Row

NAME_QUERY = '''
    select firstname, lastname
    from presidents
    where termnum < 5
'''

cur = s3conn.cursor()

# select first name, last name from all presidents
cur.execute(NAME_QUERY)

for row in cur.fetchall():
    print(row)
print('-' * 50)

dcur = s3conn.cursor() ①

# make _this_ cursor a dictionary cursor
dcur.row_factory = sqlite3.Row ②

# select first name, last name from all presidents
dcur.execute(NAME_QUERY)

for row in dcur.fetchall():
    print(row['firstname'], row['lastname']) ③

print('-' * 50)
```

- ① default cursor returns tuple for each row
- ② Row object is tuple/dict hybrid; can be indexed by position OR column name
- ③ selecting by column name

db_sqlite_extras.py

```
('George', 'Washington')  
('John', 'Adams')  
('Thomas', 'Jefferson')  
('James', 'Madison')
```

```
-----  
George Washington  
John Adams  
Thomas Jefferson  
James Madison  
-----
```

Metadata

- `cursor.description` returns tuple of tuples
- Fields
 - `name`
 - `type_code`
 - `display_size`
 - `internal_size`
 - `precision`
 - `scale`
 - `null_ok`

Once a query has been executed, the cursor's `description` attribute is a tuple with metadata about the columns in the query. It contains one tuple for each column in the query, containing 7 values describing the column.

For instance, to get the names of the columns, you could say:

```
names = [ d[0] for d in cursor.description ]
```

For non-query statements, `cursor.description` returns `None`.

The names are based on the query (with possible aliases), and not necessarily on the names in the table.

Example

NOTE

Many database modules, including pymysql, have a dictionary cursor built in — this is just for an example you could use with any DB API module that does not have this capability. The example uses the metadata from the cursor to get the column names, and forms a dictionary by zipping the column names with the column values. Another approach would be to use a named tuple. __

db_sqlite_emulate_dict_cursor.py

```
#!/usr/bin/env python

import sqlite3

s3conn = sqlite3.connect("../DATA/presidents.db")

c = s3conn.cursor()

def row_as_dict(cursor):
    '''Generate rows as dictionaries'''
    column_names = [desc[0] for desc in cursor.description]
    for cursor_row in cursor.fetchall():
        row_dict = dict(zip(column_names, cursor_row))
        yield row_dict

# select first name, last name from all presidents
num_recs = c.execute('''
    select lastname, firstname
    from presidents
''')

for row in row_as_dict(c):
    print(row['firstname'], row['lastname'])
```

db_sqlite_emulate_dict_cursor.py

```
Richard Milhous Nixon  
Gerald Rudolph Ford  
James Earl 'Jimmy' Carter  
Ronald Wilson Reagan  
George Herbert Walker Bush  
William Jefferson 'Bill' Clinton  
George Walker Bush  
Barack Hussein Obama  
Donald J Trump  
Joseph Robinette Biden
```

See `db_sqlite_named_tuple_cursor.py` for a similar example that creates named tuples rather than dictionaries for each row.

Transactions

- Transactions allow safer control of updates
- `commit()` to save transactions
- `rollback()` to discard
- Default is autocommit off
- `autocommit=True` to turn on

Sometimes a database task involves more than one change to your database (i.e., more than one SQL statement). You don't want the first SQL statement to succeed and the second to fail; this would leave your database in a corrupt state.

To be certain of data integrity, use **transactions**. This lets you make multiple changes to your database and only commit the changes if all the SQL statements were successful. For all packages using the Python DB API, a transaction is started when you connect. At any point, you can call `CONNECTION.commit()` to save the changes, or `CONNECTION.rollback()` to discard the changes. For most packages, if you don't call `commit()` after modify a table, the data will not be saved.

NOTE You can also turn on autocommit, which calls `commit()` after every statement.

Example

```
try:
    for info in list_of_tuples:
        cursor.execute(query,info)
except SQLError:
    dbconn.rollback()
else:
    dbconn.commit()
```

NOTE **pymysql** only supports transaction processing when using the **InnoDB** engine

Object-relational Mappers

- No SQL required
- Maps a class to a table
- All DB work is done by manipulating objects
- Most popular Python ORMs
 - SQLAlchemy
 - Django (which is a complete web framework)

An Object-relational mapper is a module or framework that creates a level of abstraction above the actual database tables and SQL queries. As the name implies, a Python class (object) is mapped to the actual table.

The two most popular Python ORMs are SQLAlchemy which is a standalone ORM, and Django ORM. Django is a comprehensive Web development framework, which provides an ORM as a subpackage. SQLAlchemy is the most fully developed package, and is the ORM used by Flask and some other Web development frameworks.

Instead of querying the database, you call a search method on an object representing a table. To add a row to the table, you create a new instance of the table class, populate it, and call a method like `save()`. You can create a large, complex database system, complete with foreign keys, composite indices, and all the other attributes near and dear to a DBA, without writing the first line of SQL.

You can use Python ORMs in two ways.

One way is to design the database with the ORM. To do this, you create a class for each table in the database, specifying the columns with predefined classes from the ORM. Then you run an ORM command which executes the queries needed to build the database. If you need to make changes, you update the class definitions, and run an ORM command to synchronize the actual DBMS to your classes.

The second way is to map tables to an existing database. You create the classes to match the schemas that have already been defined in the database. Both SQLAlchemy and the Django ORM have tools to automate this process.

NoSQL

- Non-relational database
- Document-oriented
- Can be hierarchical (nested)
- Examples
 - MongoDB
 - Cassandra
 - Redis

A current trend in data storage are called "NoSQL" or non-relational databases. These databases consist of *documents*, which are indexed, and may contain nested data.

NoSQL databases don't contain tables, and do not have relations.

While relational databases are great for tabular data, they are not as good a fit for nested data. Geo-spatial, engineering diagrams, and molecular modeling can have very complex structures. It is possible to shoehorn such data into a relational database, but a NoSQL database might work much better. Another advantage of NoSQL is that it can adapt to changing data structures, without having to rebuild tables if columns are added, deleted, or modified.

Some of the most common NoSQL database systems are MongoDB, Cassandra and Redis.

Example

mongodb_example.py

```
#!/usr/bin/env python
import re
from pymongo import MongoClient, errors

FIELD_NAMES = (
    'termnumber lastname firstname '
    'birthdate '
    'deathdate birthplace birthstate '
    'termstartdate '
    'termenddate '
    'party'
).split() ①

mc = MongoClient() ②

try:
    mc.drop_database("presidents") ③
except errors.PyMongoError as err:
    print(err)

db = mc["presidents"] ④

coll = db.presidents ⑤

with open('../DATA/presidents.txt') as presidents_in: ⑥
    for line in presidents_in:
        flds = line[:-1].split(':')
        kvpairs = zip(FIELD_NAMES, flds)
        record_dict = dict(kvpairs)
        coll.insert_one(record_dict) ⑦

print(db.list_collection_names()) ⑧
print()

abe = coll.find_one({'termnumber': '16'}) ⑨
print(abe, '\n')

for field in FIELD_NAMES:
    print("{0:15s} {1}".format(field.upper(), abe[field])) ⑩

print('-' * 50)

for president in coll.find(): ⑪
    print("{0[firstname]:25s} {0[lastname]:30s}".format(president))
```

```

print('-' * 50)

rx_lastname = re.compile('^roo', re.IGNORECASE)
for president in coll.find({'lastname': rx_lastname}): ⑫
    print("{0[firstname]:25s} {0[lastname]:30s}".format(president))
print('-' * 50)

for president in coll.find({"birthstate": 'Virginia'}): ⑬
    print("{0[firstname]:25s} {0[lastname]:30s}".format(president))

print('-' * 50)
print("removing Millard Fillmore")
result = coll.delete_one({'lastname': 'Fillmore'}) ⑭
print(result)
result = coll.delete_one({'lastname': 'Roosevelt'}) ⑭
print(result)
print('-' * 50)

result = coll.delete_one({'lastname': 'Bush'})
print(dir(result))
print()

result = coll.count_documents({}) ⑮
print(result)

for president in coll.find(): ⑪
    print("{0[firstname]:25s} {0[lastname]:30s}".format(president))
print('-' * 50)

animals = db.animals

print(animals, '\n')

animals.insert_one({'name': 'wombat', 'country': 'Australia'})
animals.insert_one({'name': 'ocelot', 'country': 'Mexico'})
animals.insert_one({'name': 'honey badger', 'country': 'Iran'})

for doc in animals.find():
    print(doc['name'])

```

- ① define some field name
- ② get a Mongo client
- ③ delete 'presidents' database if it exists
- ④ create a new database named 'presidents'

- ⑤ get the collection from presidents db
- ⑥ open a data file
- ⑦ insert a record into collection
- ⑧ get list of collections
- ⑨ search collection for doc where termnumber == 16
- ⑩ print all fields for one record
- ⑪ loop through all records in collection
- ⑫ find record using regular expression
- ⑬ find record searching multiple fields
- ⑭ delete record
- ⑮ get count of records

mongodb_example.py

William Howard	Taft
Woodrow	Wilson
Warren Gamaliel	Harding
Calvin	Coolidge
Herbert Clark	Hoover
Franklin Delano	Roosevelt
Harry S.	Truman
Dwight David	Eisenhower
John Fitzgerald	Kennedy
Lyndon Baines	Johnson
Richard Milhous	Nixon
Gerald Rudolph	Ford
James Earl 'Jimmy'	Carter
Ronald Wilson	Reagan
William Jefferson 'Bill'	Clinton
George Walker	Bush
Barack Hussein	Obama
Donald John	Trump
Joseph Robinette	Biden

Collection(Database(MongoClient(host=['localhost:27017'], document_class=dict,
tz_aware=False, connect=True), 'presidents'), 'animals')

wombat
ocelot
honey badger

Chapter 1 Exercises

Exercise 1-1 (president_sqlite.py)

For this exercise, you can use the SQLite3 database provided, or use your own DBMS. The `mkpres.sql` script is generic and should work with any DBMS to create and populate the presidents table. The SQLite3 database is named **presidents.db** and is located in the DATA folder of the student files.

The data has the following layout

Table 5. Layout of President Table

Field Name	Data Type	Null	Default
termnum	int(11)	YES	NULL
lastname	varchar(32)	YES	NULL
firstname	varchar(64)	YES	NULL
termstart	date	YES	NULL
termend	date	YES	NULL
birthplace	varchar(128)	YES	NULL
birthstate	varchar(32)	YES	NULL
birthdate	date	YES	NULL
deathdate	date	YES	NULL
party	varchar(32)	YES	NULL

Refactor the **president.py** module to get its data from this table, rather than from a file. Re-run your previous scripts that used `president.py`; now they should get their data from the database, rather than from the flat file.

NOTE

If you created a `president.py` module as part of an earlier lab, use that. Otherwise, use the supplied `president.py` module in the top folder of the student files.

Exercise 1-2 (add_pres_sqlite.py)

Add the next president to the presidents database. Just make up the data — let's keep this non-political. Don't use any real-life people.

SQL syntax for adding a record is

```
INSERT INTO table ("COL1-NAME",...) VALUES ("VALUE1",...)
```

To do a parameterized insert (the right way!):

```
INSERT INTO table ("COL1-NAME",...) VALUES (%s,%s,...) # MySQL
INSERT INTO table ("COL1-NAME",...) VALUES (?,?,...)   # SQLite
```

or whatever your database uses as placeholders

NOTE | There are also MySQL versions of the answers.

Index

A

API, 2
autocommit, 26

C

Cassandra, 28
commit, 26
connect(), 3
context manager, 3
cursor, 5
cursor.description, 23
cx_oracle, 2

D

database object, 5
database programming, 2
database server, 3
DB API, 2
dictionary cursor, 19
 emulating, 24
Django, 27
Django ORM, 27

E

execute(), 12, 6
executemany(), 12
executing SQL statements, 6

F

fetch methods, 7
Firebird (and Interbase, 2

I

IBM DB2, 2
Informix, 2
informixdb, 2
ingmod, 2
Ingres, 2

K

KInterbasDB, 2

M

metadata, 23
Microsoft SQL Server, 2
MongoDB, 28
MySQL, 2

N

non-query statement, 12
non-relational, 28
NoSQL, 28

O

Object-relational mapper, 27
ODBC, 2
Oracle, 2
ORM, 27

P

parameterized SQL statements, 12
placeholder, 12
PostgreSQL, 2
psycopg2, 2
PyDB2, 2
pymssql, 2
pymysql, 2
pyodbc, 2

R

Redis, 28
rollback, 26

S

SAP DB, 2
sapdbapi, 2
SQL code, 5
SQL data integrity, 26
SQL injection, 10
SQL queries, 6
SQL statements, 6
SQLAlchemy, 27
SQLite, 2
sqlite3, 2

Sybase, 2

T

transactions, 26